

Approximate Matching

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

COMP90049 Knowledge Technologies

Justin Zobel and Rao Kotagiri (CIS)

Semester 1



THE UNIVERSITY OF

MELBOURNE

Closest match

The task

Search with error

Spelling
correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic
matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of
effectiveness

Effectiveness

Relevance

Measures

Experiments

Closest match



A common computational task is to find the nearest match to a given string:



- ▶ Spelling correction, name matching, query repair, phonetic matching, data cleansing, genomics.

Further challenge: find strings with the nearest match to any of their substrings. (Why is this harder?)

Methods that address this task are “knowledge technologies” as defined in this subject.

- ▶ Thorough or exhaustive solutions (in some contexts) require unrealistic computational resources, so we need to find workable approximations.
- ▶ Correctness isn't well-defined.



Another perspective: this is *search with uncertainty*.



Search with uncertainty

Consider a file of names such as

K. Many Chandy, Jayadev Misra
Susanne Graf, Joseph Sifakis
D. Harel, A. Pnueli
Koichi Furukawa, Katsumi Niita, Yuji Matsumoto
Z. Farkas, P. Szeredi, E. Santane-Toth
Agneta Eriksson, Anna-Lena Johansson, Sten-Ake Tarnlund
Luis Moniz Pereira, Antonio Porto
Lynette Hirschmann, Karl Puder
Yuji Matsumoto, Hozumi Tanaka, Masaki Kiyono
J. Darlington, A.J. Field, H. Pull
Masahiko Sato, Tatakfumi Sakurai
P.A. Subrahmanyam, Jia-Huai You
Joseph A. Goguen, Jose Meseguer
Yonathan Malachi, Zohar Manna, Richard Waldinger



How to find names in this file, if the precise spelling is unknown?

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Search with uncertainty

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Again, **context matters**: not just **language and alphabet**, but the reason that the search is taking place.

In designing or choosing a **search method**, three **factors need** to be considered:

- ▶ **What** the user is **trying to achieve**.
- ▶ The **source of reference** material for “**correct strings**”.
- ▶ **Tractability** and **interactivity** – is a user waiting for a response? How will the solution behave as the uncertainty increases?

Tools: ispell, agrep, web search interfaces.

Example: the **agrep** tool is an implementation of an algorithm for matching with errors.

Pattern matching tools and spelling correction

Closest match

The task
Search with error

Spelling
correction

As pattern matching

Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

The motivating use of pattern matching tools such as `agrep` is to identify likely matching lines in a file. No assumptions are made about meaning of characters.

The aim of spelling correction is to:

- ▶ Parse text into “words”. (We’ll look at parsing in another lecture.)
- ▶ Ignore the words that are found in a dictionary.
- ▶ Find the best matches for the remainder.

... proceed if you can see no `ther` option ...

Matches with `agrep -1 ther: %` with one correction
ether, her, other, thar, the, thee, their, them, then, there,
therm, thew, they, tier



Observe the behaviour of `agrep` as k (number of corrections: insertions, deletions or substitutions) is increased.

Pattern matching tools and spelling correction

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments



A straightforward process, but it does expose underlying concepts:

- *Efficiency* and *ranking*.

Ranking. “Best match” is relative. Whether a string is a match to a query is not a binary (yes/no) decision, and the goal should be to present matches according to perceived *likelihood* of their being correct.

Note that the task inherently requires *human participation*.

What other tasks use *ranking*?



Pattern matching tools and spelling correction

Closest match

The task
Search with error

Spelling correction

As pattern matching

Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Efficiency. Exhaustive file search is $O(n)$ in time, for say a dictionary of n entries.

Most words should be found in the dictionary; there are plenty of data structures that provide effective $O(1)$ confirmation of whether a particular string is an exact match (e.g., hashing or tries).

(Yes, I do mean *tries*, not *trees*.)

Ballpark: 10^6 to 10^7 searches per second on a laptop.

$O(n)$ may be unsatisfactory for the approximate search task. Perhaps more importantly, processing a file includes overheads such as parsing text into lines – doing so for each word to be checked is a waste of resources.

Pattern matching tools aren't customized for parsing; they are not efficient at eliminating exact matches; and don't rank results.

Conclusion: Despite the task similarity, they are not suitable for spell checking.

Aside: stemming

There are two kinds of dictionaries used in spelling correction.

- ▶ **Literal:** a set of strings that must be matched exactly.
- ▶ **Root+stem:** a set of root words, a set of suffixes, and a set of rules for applying suffixes to roots.

Root: fill, apply, invoke.

Word:

filled, filling, but not fillcation

applied, application, but not apple

invoked, invoking, invocation, but not invokery

Stemming is the process of stripping suffixes automatically. It allows a dictionary to be generalized, but many root-suffix combinations are not correct spellings.



Observe that such technologies are highly language-dependent.

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Closest match

The task
Search with errorSpelling
correction

As pattern matching

Stemming

String
neighbourhood

Edit distance

N-grams
Query correctionPhonetic
matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of
effectiveness

Effectiveness

Relevance

Measures

Experiments

Neighbourhood search

Given a query string that is not in the dictionary, the task is to find the nearest neighbours.

One approach is to generate all neighbours, and see if they are present.

- ▶ Enumerate all of the single-character deletions, insertions, and replacements.

This generates all strings whose distance from the query is 1.

For example, for an alphabet of four letters {e h r t}, the word **thr** has the immediate neighbours {hr tr th ethr hthr rthr tthr tehr thhr trhr tthr ther thhr thrr thtr thre thrh thrr thrt ehr hhr rhr ter trr ttr the thh tht}

How many neighbours? Alphabet size A , query length n :

$$N_1 = n + A \cdot (n + 1) + (A - 1) \cdot n$$

n is due to one character deletions; $A \cdot (n + 1)$ is due to one character insertions; $(A - 1) \cdot n$ is due to one character replacements; Ballpark: $A = 26$, $n = 8$, so $N_1 = 441$ and $\gg 1000$ unknown strings can be processed per second.

Looks feasible! This is how `ispell` works.

Costs

The cost of neighbourhood search is $O(n)$ since the alphabet is fixed.

What happens if we want to increase the distance from the query?
For example, distance of at most 2:

$$N_2 = n + n \cdot (n - 1) + A \cdot (n + 1) + A \cdot (n + 1)A \cdot (n + 2) + \dots$$

This does not include combinations of deletion, replacement, and insertion. n is due to one character deletions; $n \cdot (n - 1)$ is due to two character deletions; $A \cdot (n + 1)$ is due to one character insertions; $A \cdot (n + 1)A \cdot (n + 2)$ is due to two character insertions;

Ballpark: $N_2 \gg 50,000$, and cost is $O(n^2)$ with a high (≈ 600) constant factor.

This combinatoric analysis indicates that exhaustive search beyond the immediate (distance = 1) neighbourhood is at the upper limit of feasibility – and usability.

Closest match

The task
Search with error

Spelling
correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

Neighbourhood size

Query is ther. What is going on here?



aer aether ahem anther aver bather be^{er}er pier bother char cheer chef
chert chew chez cither daer deer dither doer dyer either etcher
ether ever ewer father fer gather ghee goer he hear heir hem hen
hep her herb herd here herl herm hern hero hers hes hew hex hey
hither hoer jeer kier lather leer lither mither mother nether oer
omer other others otter over oyer peer per phew pier pother rather
rhea seer she shear shed sheer sherd shes shew shier shoer shyer
steer suer taed tael tahr taker taler tamer taper tar tater taxer
tea tear tee teed teem teen tees ten term tern tether thaler than
thar that thaw the thebe theca thee theed thees theft thegn their
theirs them theme then thenar theory there therm therms these theta
thew thews thewy they thief thin third this tho thorn thorp thou
three threw thru thud thug thus thy tie tied tier tiers ties tiger
tiler timer tither toe toea toed toes toner toper tor torr toter
tother tour tower tree tref trek tret Trey trier truer tsar tuber
tuner twee twerp tzar user usher utter veer weer wether whee when
where whet whew whey whir wither yer zither

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Combinatoric search

Current practical techniques use variants of binary search, where:

- ▶ Strings are sorted, that is, grouped by prefix.
- ▶ Strings are inspected character-by-character until it is clear that their distance from the query is greater than k .
- ▶ Due to the sorting, groups of strings that share a prefix can be considered simultaneously.

Since the string space is extremely sparse (of the possible A^n strings, only a tiny fraction are actually observed), and there are good special-purpose data structures (e.g., suffix arrays), the cost in practice is around that of neighbourhood search at distance 1.

The distance limit k must be chosen, or repeated search is required. Additional work is required for ranking.

Because there is no sorting of a dictionary that ensures that similar strings are close together, it is still not obvious what the best solution to this search problem is.

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Aside: Spelling correction in context

Closest match

The task

Search with error

Spelling
correction

As pattern matching

Stemming

String
neighbourhood

Edit distance

N-grams

Query correction

Phonetic
matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of
effectiveness

Effectiveness

Relevance

Measures

Experiments

In principle, it might seem appealing to use grammar or morphological structure to help guide spelling correction.

But there are significant obstacles:

- ▶ Relatively few spelling errors can be resolved by context.
- ▶ The cost of building the dictionary would be prohibitive.

There are good contextual spelling checkers for typical errors, based on both machine learning and manual analysis of syntax and error patterns. These are most effective for short words.

There are no spelling checkers that use broad context, such as the category or topic of a document (however this might be defined). We are probably stuck with order-0 (context-free) correction for most words. But we can try and be smarter about what we try to correct ...

Edit distances

From another perspective: take the query string and each candidate match, and find the *edit distance* between them.

The simplest edit distance is the number of insertions, deletions, and substitutions needed to turn one string into another.

The smaller the edit distance, the closer the match.

Gorbach—ev

 | |
Gorbechyov

Connell

 |
Con—al—

Measures such as edit distances are attractive because matches are scored, so answers can be ranked.

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Edit distances

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String
neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Consider the alignment of two strings `agt` and `aact` using a simple scheme where an alignment of two matching characters scores $+1$, an alignment of two mismatching characters scores -1 , and an insertion or deletion of a character also scores -1 .

(There is no explicit penalty for sequences being of unequal lengths).

Intuitively, using this simple scoring scheme, an optimal alignment of the two sequences would be:
$$\begin{bmatrix} \text{aact} \\ -\text{agt} \end{bmatrix}$$



Edit distances

Assume an array F of size $|q| \times |t|$, to be used to compute the *global* edit distance between q and t .

NEEDLEMAN-WUNSH Algorithm:

```
lq = strlen(q); lt = strlen(t);
for( i=0 ; i<=lq ; i++ ) F[i][0] = -i;
for( j=0 ; j<=lt ; j++ ) F[0][j] = -j;

for( i=1 ; i<=lq ; i++ )
    for( j=1 ; j<=lt ; j++ )
        F[i][j] = max3(
            F[i-1][j] -1, % insertion
            F[i][j-1] -1, % deletion
            % match/miss match
            F[i-1][j-1] + equal(q[i-1], t[j-1]));
```

`equal()` returns +1 if equal, -1 otherwise.

The value $F[lq][lt]$ is the minimum number of edits between q and t .
(But what is an edit?)

This method is known for historical reasons as dynamic programming.

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood

Edit distance

N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Edit distances

The *local* edit distance is similar.
SMITH-WATERMAN algorithm:

```
lq = strlen(q); lt = strlen(t);
for( i=0 ; i<=lq ; i++ ) F[i][0] = 0;
for( j=0 ; j<=lt ; j++ ) F[0][j] = 0;

for( i=1 ; i<=lq ; i++ )
    for( j=1 ; j<=lt ; j++ )
        F[i][j] = max4(
            0,
            F[i-1][j] - 1, % insertion
            F[i][j-1] - 1, % deletion
            match/miss match
            F[i-1][j-1] + equal(q[i-1], t[j-1])
        );
```

Here, `equal()` returns `+1` if equal, `-1` otherwise. (But there are many choices of scoring schemes, as we discuss later.)

The value of `F[i][j]` with the *highest* value represents the best alignment.

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Edit distances

There are three possibilities in the alignment of the first two characters in each sequence: $\begin{bmatrix} a \\ a \end{bmatrix}$, $\begin{bmatrix} a \\ - \end{bmatrix}$, and $\begin{bmatrix} - \\ a \end{bmatrix}$.

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics


Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments



Edit distances

Since $+1$, using alignment $\begin{bmatrix} a \\ a \end{bmatrix}$, is the maximum score, we store the score $+1$ in $[1, 1]$ and mark the path used to calculate the optimal score.

Closest match

The task
Search with error

Spelling correction

As pattern matching

Stemming
String
neighbourhood

Edit distance

N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

		a	a	c	t
a	0	0			
g	0	1			
t					

Closest match

The task

Search with error

Spelling
correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic
matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of
effectiveness

Effectiveness

Relevance

Measures

Experiments

Having considered the first cell, we consider all other cells in the matrix.

		a	a	c	t
		0	0	0	0
a		0	1	1	0
			1	0	-1
g		0	0	0	0
			0	0	-1
t		0	-1	-1	1
			-1	-1	1

Edit distances

In some cases, such as that of cell $[4, 1]$, the maximum score can be achieved through more than one alignment; in this case, since the scores of cells $[4, 0]$, $[3, 0]$ and $[3, 1]$ are all 0, all paths contribute to a total score in $[4, 1]$ of -1 .

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

	a	a	c	t
a			0	0
g			0	-1
t				

Edit distances

The alignment contributing to the maximum score can be reconstructed by tracing-back the pointers stored during the calculation. In this case, there is only one path from [4, 3] that is highlighted with a dashed

ellipse: $\begin{bmatrix} \text{aact} \\ -\text{agt} \end{bmatrix}$

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood

Edit distance


N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments



Cost of edit distance

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Global alignment of strings was first proposed by Needleman and Wunsch (1970), but is based on at least seven published schemes.

The original formulation of Needleman and Wunsch was $O(n^3)$ in time and $O(n^2)$ in space; their approach considered each cell in row i and each cell in row j in calculating a value for $[i, j]$.

The approach described here is $O(n^2)$ in both time and space.

For short strings, edit distance is practical. For longer strings or comparison against a database of strings, the time cost may be unreasonable.

It isn't obvious how to find strings to test; exhaustive testing is probably infeasible.

N-grams

Let $G_n(s)$ be the substrings of length n of s .

Example1: $G_2(\text{Gorbachev})$ is Go, or, rb, ba, ac, ch, he, ev

Example2: $G_2(\text{Gorbechyov})$ is Go, or, rb, be, ec, ch, hy, yo, ov

The n-gram distance of t and s is

$$|G_n(s)| + |G_n(t)| - 2 \times |G_n(s) \cap G_n(t)|$$

Example: $s = \text{Gorbachev}$, $t = \text{Gorbechyov}$, distance is
 $8 + 9 - 2 \times 4 = 9$.

The smaller the n-gram distance, the closer the match. Cost is approximately $O(|s| + |t|)$, compared to $O(|s| \cdot |t|)$ for an edit distance.

A useful variant, but not much cheaper than an edit distance in practice – it still requires an exhaustive search of the candidates.

Seems unlikely to produce useful scores for long strings or small alphabets.

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance

N-grams

Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Aside: Query correction

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Web search tools offer corrections to query spellings, typically one or zero alternatives per query.

What is happening under the hood?

Hint: what might the reference material be?

Phonetic matching

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching

Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Many databases include vocabularies of words or names:

- ▶ Surnames in the telephone directory.
- ▶ Distinct words occurring in a text database.
- ▶ Place names in a gazette.

Phonetic matching techniques allow searching for a name or word when the exact spelling is unknown:

- ▶ A user may only know a name by its sound.
- ▶ A name can be misspelt in a database.
- ▶ Some names have variant or indeterminate spelling.

Example

Georgia Conal

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Example

Georgia Conal
George O'Connell

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Example

Georgia Conal
George O'Connell

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Spelling of names is highly variable:

Lho, Lo, Loan, Loe, Loew, Lough, Low, Lowe, ...

In ≈ 1 Gb of newspaper articles the name “Gorbachev” had > 20 variants including:

*Gorbachev, Gorbacahev, Gorbahev, Gorbachev, Gorbechev,
Gorbachov, Gorachev, Gorbacheva, Gorbechyev, Gorbacev,
Gorbachyov, Gorabchev, Grobachev, ...*

Why?

Example

Georgia Conal
George O'Connell

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Spelling of names is highly variable:

Lho, Lo, Loan, Loe, Loew, Lough, Low, Lowe, ...

In ≈ 1 Gb of newspaper articles the name “Gorbachev” had > 20 variants including:

*Gorbachev, Gorbacahev, Gorbahev, Gorbachev, Gorbechev,
Gorbachov, Gorachev, Gorbacheva, Gorbechyev, Gorbacev,
Gorbachyov, Gorabchev, Grobachev, ...*

Why?

Simple case: Schröder \rightarrow Schroder or Schroeder?

Soundex

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String
neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

SoundexLexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

The Soundex sound-alike matching technique is well-known, widely used, simple to implement.

Transforms a string into a 4-character code that represents its sound:

- ▶ Replace all but the first letter by one of seven single-digit codes:

a e h i o u w y → 0 b f p v → 1

c g j k q s x z → 2 d t → 3

l → 4 m n → 5 r → 6

- ▶ Remove doubles, then remove 0s.
- ▶ Truncate to four symbols.

Soundex examples

a e h i o u w y \rightarrow 0 b f p v \rightarrow 1
c g j k q s x z \rightarrow 2 d t \rightarrow 3
l \rightarrow 4 m n \rightarrow 5 r \rightarrow 6

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

King Kyngge

⇒ K052 K05220

⇒ K052 K0520

⇒ K52 K52

Knight Night

⇒ K50203 N0203

⇒ K50203 N0203

⇒ K523 N23

Loan Loew Lough Lewicks

⇒ L005 L000 1002 L000222

$$\Rightarrow \quad L05 \quad L0 \quad L02 \quad L02$$
$$\Rightarrow \quad L5 \quad L \quad L2 \quad L2$$

There are 6734 ($26 \times (1 + 6 + 6^2 + 6^3)$) distinct Soundex codes.

There are also (failed) variants, such as Phonix.

Lexicographic methods

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Strings that sound alike are often spelt alike.

Example: Gorbachev, Gorbechyov, ...

It is plausible to suppose that lexicographic string comparison methods will be effective at finding strings that sound alike.

Candidates include neighbourhood search, edit-distances, and n-gram measures.

Perhaps these would be enhanced by combining them with approaches that embody language knowledge.

Refining edit distances

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Based on lexicographic information alone, matches such as “Crews” and “Kroose” would not be highly ranked.

Can edit distances be modified to allow for phonetic similarity somehow?

- ▶ Repeat letters don't usually change the sound;
“Conel” sounds like “Connell”.
- ▶ Sometimes different letters give rise to the same sound;
“Kodi” sounds like “Cody”.
- ▶ But sometimes they don't;
“like” doesn't sound like “lice”.

“Editex”

Exploring some possibilities ... can edit distances be phonetically tweaked?

Zobel’s **editex**: Like the standard edit distance, but instead of having a fixed penalty for insertion, deletion, and substitution, have two penalties:

- ▶ **High** for letters that are **never similar**, such as “d” and “m”.
- ▶ **Low** for letters that can give rise to similar sound, such as “m”–“n” and “c”–“k”.

Ten	1. a e i o u y	2. b p	3. c k q	4. d t	5. l r
groupings:	6. f p v	7. s x z	8. c s z	9. m n	10. g j

(The silent letters “w” and “h” are handled as a special case.)

Crews—

ʔ ʔ |

Kroose

Closest match

The task
Search with error

Spelling
correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

But what about phonetics?

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Letters don't represent sounds: consider "t" in "trim" versus "think".

"th" represents a sound: the phoneme "θ" in the International Phonetic Alphabet.

Every name has a pronunciation, or, more precisely, can be represented as a string of phonemes. Perhaps these can be used for matching in some way.

Two problems:

- ▶ Deciding what pronunciation corresponds to a given string.
- ▶ Comparing pronunciations.

Closest match

The task
Search with error

Spelling
correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

- ▶ Use a text-to-sound algorithm to generate the likely pronunciation of each string.
- ▶ Use some form of edit distance to compare pronunciations.

Crews	→	krjəs
Kroose	→	kr—əs

A more principled solution but:

- ▶ Text-to-sound algorithms are highly reliable only if context is available, for example to distinguish “in a minute” from “it is minute”.
- ▶ The pronunciation of names is much less predictable than the pronunciation of words.

The following are the Physiological Symbols for the English elements of Speech.

O p in pea.	u t in tea.	a k in key.	o r in train.
o b in bay.	o d in day.	e g in gay.	o r in rain.
o m in some.	o n in son.	e ng in sung.	o h in hue.
o f in fine.	o th in thigh.	o l in cloud.	o y in you.
o v in vie.	o th in thy.	o l in loud.	o h in hop.
o wh in whey.	o s in hiss.	o sh in rush.
o w in way.	o s in his.	o ge in rouge.

f ee in eel.	f i in ill.	ſ e in shell.	ſ a in shall.
þ oo in pool.	þ u in pull.	þ a in all.	þ o in doll.
ſ a in father.	ſ a in ask.	ſ u in curl.	ſ u in dull.

4 w as in now. y r as in sir. ay as in may. I a as in near.

i in mine. [a in mane. ow in now. ow in know.
oy in boy.

Illustration of the Physiological Alphabet.

සහ අන්තර්ජාතික පදනම විසින් 1876 වර්ෂයේ දී
 The Commissioners of the International Exhibition of 1876 have granted
 මෙම ත්‍යාගය ජිව විද්‍යාත්මක අක්ෂරලේඛනයක් සඳහා
 an award for the Physiological Alphabet devised by Professor A.
 බ්‍රැන්ට්ස්ටන් විසින්
 Melville · Bell, of Brantford, Ontario.

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String
neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Measurement of efficiency is straightforward.

But we also need to know whether the method is useful. Measurement of *effectiveness* involves human assessment of the quality of the results.

For each task we need to find an independent method that measures effectiveness in a reliable way.

- ▶ Choose a test lexicon and a set of query strings.
- ▶ For each query and each word in the lexicon, use relevance assessors to decide whether the word and query match, that is, the word is *relevant* to the query.
- ▶ Apply the matching technique to the lexicon and each query to get a list of possible answers, and measure effectiveness by examining how many of the answers are relevant.

Relevance

Applying `agrep -1 davis` to a file of names drawn from internet news postings:

avis, danis, david, davie, davies, davis,
daviss, davys, lavis, navis, ravis

Using a `human to judge` which of these matches are `probably correct`, also considering some other matching techniques and some other names, yields:

barlow	bulow balo barlow
bloom	bloom blume bluhm
clark	klerk kluch clack clerc clarke clark
davis	daviss davyes davys davis
farah	faraj farah farra farrall farrar
hahn	hahm hann hahn hanne

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Using relevance information

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Given a query, a matching technique A returns a ranked list of answers.

Given relevance judgements, a top-20 ranking from A might be represented as

◇ ◇ ○ ◇ ○ | ○ ○ ◇ ◇ ○ | ○ ◇ ? ◇ ○ | ◇ ○ ? ? ○

where ◇ is “correct”, ○ is “incorrect”, and ? is unjudged. Assuming that unjudged matches are incorrect, this ranking is equivalent to

◇ ◇ ○ ◇ ○ | ○ ○ ◇ ◇ ○ | ○ ◇ ○ ◇ ○ | ◇ ○ ○ ○ ○

Using relevance information

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Reminder: A 's ranking is:

◇ ◇ ○ ◇ ○ | ○ ○ ◇ ◇ ○ | ○ ◇ ○ ◇ ○ | ◇ ○ ○ ○ ○

Suppose that method B returns the ranking

○ ◇ ◇ ◇ ◇ | ○ ○ ○ ◇ ○ | ○ ○ ◇ ○ ○ | ○ ◇ ○ ◇ ◇

To decide which method has the **greater effectiveness**, we need to agree on a method for scoring a ranking.

There is **no objective criteria** that can be used to **assess a scoring method**! Choice of measure is based on arguments from properties such as **expected user behaviour**.

Using relevance information

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

The traditional measures of effectiveness are **precision** and **recall**.

Precision: The proportion of the retrieved matches **that are correct**.

Reminder – the rankings were:

◇ ◇ ○ ◇ ○ | ○ ○ ◇ ◇ ○ | ○ ◇ ○ ◇ ○ | ◇ ○ ○ ○ ○
○ ◇ ◇ ◇ ◇ | ○ ○ ○ ◇ ○ | ○ ○ ◇ ○ ○ | ○ ◇ ○ ◇ ◇

In the top 20, **A has precision 40%** (8 out of 20) and **B has precision 45%** (9 out of 20).

The limit of 20 was an arbitrary cut-off; we might instead say, for example, **that P@10 for both A and B is 50%**.

Using relevance information

Closest match

The task
Search with error



Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Recall: The proportion of the correct matches that are retrieved.

Reminder – the rankings were:



We do not know how many correct matches there are for this query. If we assume that 14 strings have been judged correct (not all correct were judged), then *A* has recall 57% ($8 \div 14$) and *B* has recall 64% ($9 \div 14$).

Using relevance information

Closest match

The task
Search with error

Spelling
correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

As users tend to proceed down a ranking examining answers, in turn, an alternative is use measures such as average precision (AP) – the average of the precisions at each correct matches retrieved. Correct matches that are not retrieved are assigned a precision of 0.0. AP for A is 0.387, and for B is 0.363.

Formally, if r_i is 1 (respectively, 0) if the i th item in a ranking is relevant (respectively, irrelevant), there are d items in the ranking, and there are in total R known relevant items, average precision is defined as

$$AP = \frac{1}{R} \sum_{i=1}^d \left(\frac{r_i}{i} \cdot \sum_{j=1}^i r_j \right)$$

A detailed discussion of this measure is in the readings.

Using relevance information

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Different performance measures are not necessarily consistent with each other.

Given the diversity of retrieval methods used by different search engines, it is easy to have $A > B$ for one query and $A < B$ for another. For this reason, we use a sample of queries (10 is too few, 50 is adequate, more is preferable) and average the per-query effectiveness.

We then need to use a statistical test to ensure that the two samples are indeed different (to some level of confidence), rather than simply different by chance.

There are many other measures – F-score, error rate, “receiver operator characteristic” ROC and “area under ROC”, discounted cumulative gain, rank-biased precision ...

The importance of baselines

We use *baselines* to establish whether any proposed method is **doing better than** “dumb and simple” – “dumb” methods often work surprisingly well.

Baselines are also valuable in getting a sense for the intrinsic difficulty of a given task – sometimes a simple method is so good that there isn’t much scope for a clever method to do better.

Example: identifying **what language** a document is written in. Just need to check which languages have the **document’s commonest words**.

In formulating a baseline, we need to be sensitive to the task and how the method is to be used.

Example:

- ▶ In **spelling correction** it is **helpful to give a lot of options; only one of them will be correct**.
- ▶ In **query correction**, an option is only shown **if there is strong evidence that it is correct**.

The baseline should be **plausible** – there is no point in comparing to the worst method available.

Baselines vs. benchmarks

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Baseline – naïve or straightforward method that we would expect a rich (or principled, or . . .) method to do better.

Benchmark – established “best practice” technique that we are pitching our method against.

The word “baseline” is often used as an umbrella term for both meanings.

Building a test collection

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

Names from the “From:” lines of Internet news articles, hand-edited to standardise format and eliminate obvious rubbish (Darth Vader, GVR 101187-9456, The American Psychological Association).

Queries selected by generating random page numbers in the range 80–1800 and using them as an index into the Melbourne White Pages.

With each of the 125 (!) phonetic matching techniques implemented, fetch the best matches and form a pool.

Using a pair of relevance assessors to judge the pooled answers to each query (one speaking, one listening).

Three teams \Rightarrow three sets of relevance judgements.

Retrieval performance

Average **precision** (in parentheses, number correct at **200**). Each column is a different speaker–assessor team.

Method	Set of judgements		
	A	B	C
Editex	23.1 (17.8)	28.5 (4.3)	26.7 (6.9)
Ipadist	23.2 (16.4)	23.2 (3.8)	24.5 (6.1)
Edit distance	20.4 (15.4)	25.1 (3.7)	22.3 (6.3)
N-gram	20.0 (16.5)	21.5 (3.5)	22.2 (6.2)
Agrep –1	16.5 (3.5)	18.7 (1.2)	18.3 (2.1)
Agrep –best	12.1 (0.8)	20.3 (0.6)	15.2 (0.8)
Soundex	9.3 (6.0)	6.9 (1.8)	9.5 (3.2)

The orderings given by the teams are nearly identical, despite the different scores achieved.

Closest match

The task
Search with error

Spelling
correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic
matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of
effectiveness

Effectiveness
Relevance
Measures
Experiments

Independence of methods

Consider the query `crews` on the two best methods:

<i>rank</i>	<i>lpadist</i>	<i>Editex</i>
1.	<code>crews</code>	<code>crews</code>
2.	<code>krewe</code>	<code>cress</code>
3.	<code>kreuser</code>	<code>clews</code>
4.	<code>crew</code>	<code>drews</code>
5.	<code>drews</code>	<code>crew</code>
6.	<code>clews</code>	<code>creps</code>
7.	<code>kruse</code>	<code>kress</code>
8.	<code>kroose</code>	<code>cross</code>

That is, measurement based on **recall** and **precision** can be similar even if the **outcomes are very different**.

As mentioned earlier, it is necessary to use a good number of queries to be confident that one method is better than another.

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

Summary

Closest match

The task
Search with error

Spelling correction

As pattern matching
Stemming
String
neighbourhood
Edit distance
N-grams
Query correction

Phonetic matching

Name matching
Soundex
Lexicographic
methods
Phonetics

Measurement of effectiveness

Effectiveness
Relevance
Measures
Experiments

- ▶ Approximate matching techniques provide search with uncertainty.
- ▶ Here we have focused on spelling correction, as an example of a simple knowledge technology where effectiveness is important.
- ▶ Most methods use some form of distance measure, or do exhaustive search within an error bound.
- ▶ For the specific case of phonetic error, special-purpose methods are superior to purely alphabetic methods.
- ▶ Testing of methods requires a test data set; test systems or methods, including a baseline; human assessment of outcomes; an effectiveness measure.
- ▶ Such testing methods underpin investigation of areas such as text search and document classification.

Readings

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments

en.wikipedia.org/wiki/Bitap_algorithm

en.wikipedia.org/wiki/Spell_checker

www.googleguide.com/spelling_corrections.html

“Learning a Spelling Error Model from Search Query Logs” (LMS).

“Rank-Biased Precision for Measurement of Retrieval Effectiveness” (LMS), sections 2 & 3.

“Phonetic String Matching: Lessons from Information Retrieval” (LMS).

Manning et al., chapters 3 & 8.

Approximate Matching

COMP90049 Knowledge Technologies

Closest match

The task

Search with error

Spelling correction

As pattern matching

Stemming

String

neighbourhood

Edit distance

N-grams

Query correction

Phonetic matching

Name matching

Soundex

Lexicographic
methods

Phonetics

Measurement of effectiveness

Effectiveness

Relevance

Measures

Experiments